



◆ CERTIFICATION & ARCHITECTURE ROADMAP

One Storage Layer. Every Workload. Master DP-700.

Everything you need to know to pass the Microsoft Certified: Fabric Data Engineer Associate exam — simplified in 30 days.

📅 30 Days Challenge

🌐 3 Equal Domains

🎯 700 Pass Score

📖 ONELAKE FOUNDATION AT A GLANCE

Core Concept

1

Single Tenant Data Lake

OneLake unifies all workspaces without data silos

0

Zero Data Copying

Direct Lake & OneLake Shortcuts eliminate sync latency

100%

DP-700 Syllabus Aligned

Aligned with official July 2026 study guide



THE CERTIFICATION

Microsoft Fabric Data Engineer

A role-based Microsoft certification for data engineers who build solutions on Fabric.

EXAM QUICK FACTS



TARGET ROLE
Data Engineer / Analytics Engineer



PASS SCORE
700 out of 1000 points



EXAM FORMAT
40–60 questions • Online or Prometric Centre



RENEWAL
Annual — Free online assessment on Microsoft Learn



CORE LANGUAGES
SQL / T-SQL • PySpark / Python • KQL



Who Should Take This Exam?

Data engineers, analytics engineers, and architects who design, build, monitor, and optimize data pipelines, lakehouses, and warehouses using Microsoft Fabric.



THE PLATFORM

One Platform. Everything.

Microsoft Fabric unifies the entire analytics stack under a single SaaS platform.



Unified Architecture

Replaces Azure Synapse, Power BI Premium, and Data Factory — all connected through OneLake.



One Copy of Data

No more silos. OneLake stores data once and makes it available to every Fabric workload simultaneously.



Fastest-Growing Platform

Microsoft Fabric is one of their fastest-growing cloud services. Demand for Fabric engineers is exploding globally.



Enterprise Security

Built-in governance, sensitivity labels, OneLake security, RLS, OLS, and audit logs — out of the box.



The Unified Data Stack

Lakehouse • Data Warehouse • Pipelines • Notebooks • Real-Time Intelligence • Power BI • OneLake



EXAM BLUEPRINT

3 Equal Domains

Each domain carries **30–35%** of the exam weight — study all three equally.

Manage	Ingest	Monitor
Implement & Manage Analytics Workspace Settings Lifecycle & Git Security Orchestration OneLake 30–35%	Ingest & Transform Data Batch Loading Streaming PySpark SQL KQL Shortcuts 30–35%	Monitor & Optimize Error Resolution Spark Tuning DW Optimization Alerts 30–35%



Study Tip


Unlike many Microsoft exams, DP-700 is perfectly balanced. Don't over-invest in one domain — spread your study time equally across all three.



THE FOUNDATION

OneLake — How It Works

One storage layer. Every Fabric workload reads from the same Delta Parquet foundation.



- Delta Parquet
- No Data Duplication
- Cross-workload Access
- Shortcuts Support



THE TOOLKIT

9 Tools for the Exam

**Lakehouse**

Files + Delta tables. SQL analytics endpoint built-in.

Ingest

**Pipelines**

Copy & Spark activities. Schedule & event triggers.

Orchestrate

**Dataflows Gen2**

Low-code Power Query ETL. 150+ connectors.

Low-Code

**Notebooks**

PySpark & SQL. Spark 3.5 + Delta support.

PySpark

**Data Warehouse**

Serverless T-SQL. Cross-DB queries.

T-SQL

**Eventhouse / KQL**

Real-time analytics with Kusto Query Language.

KQL

**Eventstream**

Capture & route real-time event streams.

Streaming

**Shortcuts**

Virtual links to ADLS, S3, GCS without copying.

Cross-Cloud

**Deploy Pipelines**

Dev → Test → Prod CI/CD via Git.

CI/CD



PREREQUISITES

3 Languages, 1 Exam

Master these three languages and you will cover the transformation core of DP-700.

01

SQL / T-SQL

COPY INTO, CTAS, cross-DB queries, RLS, Dynamic Data Masking in the Fabric Data Warehouse workload.

Critical

02

PySpark / Python

ETL transformations, Delta table operations, structured streaming, and Spark performance optimization in Notebooks.

Critical

03

KQL — Kusto Query Language

Real-time event queries, time-series analysis, windowing functions in Eventhouse / KQL Databases.

Important

04

Data Architecture

Medallion patterns, dimensional modeling, incremental loading, streaming pipeline design, and OneLake topology.

High



THE PLAN

5 Weeks, 3 Domains



DAYS 1-6 • ARCHITECTURE & WORKSPACES

OneLake, Spark Settings, Delta Lake, Airflow

OneLake

Spark Settings

Delta Lake

Apache Airflow



DAYS 7-13 • BATCH INGESTION & TRANSFORMATION

Pipelines, Shortcuts, Mirroring, PySpark ETL

Pipelines

Shortcuts

Mirroring

PySpark ETL



DAYS 14-19 • DATA WAREHOUSING & T-SQL

DW Load Patterns, Cross-DB, RLS, DDM

COPY INTO

Cross-DB

RLS / DDM



DAYS 20-25 • REAL-TIME INTELLIGENCE & KQL

Eventstream, KQL DB, Structured Streaming

Eventstream

KQL DB

Spark Streaming



DAYS 26-30 • MONITOR, OPTIMIZE & CI/CD

Git, DB Projects, Deployment Pipelines, Spark Tuning

Git

DB Projects

Deployment Pipelines

Spark Tuning



TIP #1 OF 30

Start with OneLake Shortcuts

Before learning pipelines, understand how **OneLake Shortcuts** let you access external data without copying it. The exam frequently tests *when to use Shortcuts vs Pipelines* — Shortcuts are often the right choice.



What is a Shortcut?

A virtual link to data in ADLS Gen2, Amazon S3, or Google Cloud Storage that appears as a folder inside your Lakehouse — no data movement.



Exam Focus

Know the difference between native tables, OneLake shortcuts, and Query Acceleration shortcuts in Real-Time Intelligence.

No Data Copy

Cross-Cloud

Exam Favorite

Cost-Efficient



DAY 1 COMPLETE

Day 1 — Done!

You now understand the **DP-700 exam blueprint**, the **OneLake architecture**, and the 9 key tools you'll master over the next 30 days.

Follow on LinkedIn for Day 2 →

Tomorrow: Spark Workspace Settings — Custom pools, runtimes & executor configs

ABOUT THE AUTHOR



Sayan Banerjee

Databricks Certified Data Engineer • Microsoft Fabric Data Engineer (DP-700) • Power BI Analyst (PL-300)

AI, BI & Analytics Consultant • Embee Software • Kolkata

